

**pip install xgboost : Install XGboost if not installed.**

## Data preprocessing

### Load datasets ¶

#### Loading train.csv to train our model and observe the shape

```
In [212]: import pandas as pd  
import numpy as np
```

```
In [185]: mercedesBenz_df = pd.read_csv("E:\\Simplilean\\ML\\ProjectSubmission\\MercedesB  
enz\\Datasets\\train.csv", sep=",")
```

```
In [153]: mercedesBenz_df.shape
```

```
Out[153]: (4209, 378)
```

#### Loading test.csv to predict our new data using XGboostRegressor model and

#### Checking shape

#### Visualizing in table using head()

#### This test data will be used to predict values after our model XGboost Regressor got trained

```
In [193]: new_df = pd.read_csv("E:\\Simplilean\\ML\\ProjectSubmission\\MercedesBenz\\Data  
sets\\test.csv", sep=",")
```

```
In [194]: new_df.shape
```

```
Out[194]: (4209, 377)
```

In [156]: `new_df.head()`

Out[156]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382
0	1	az	v	n	f	d	t	a	w	0	...	0	0	0	1	0	0	0
1	2	t	b	ai	a	d	b	g	y	0	...	0	0	1	0	0	0	0
2	3	az	v	as	f	d	a	j	j	0	...	0	0	0	1	0	0	0
3	4	az	l	n	f	d	z	l	n	0	...	0	0	0	1	0	0	0
4	5	w	s	as	c	d	y	i	m	0	...	1	0	0	0	0	0	0

5 rows × 377 columns



## Splitting train data into X, y to train our model.

### printing values using head() to make confirmatiom

In [186]: `y = mercedesBenz_df.iloc[:,1:2]`

In [187]: `mercedesBenz_df.pop('y')`  
`X= mercedesBenz_df`

In [188]: `X.head(10)`

Out[188]:

	ID	X0	X1	X2	X3	X4	X5	X6	X8	X10	...	X375	X376	X377	X378	X379	X380	X382
0	0	k	v	at	a	d	u	j	o	0	...	0	0	1	0	0	0	0
1	6	k	t	av	e	d	y	l	o	0	...	1	0	0	0	0	0	0
2	7	az	w	n	c	d	x	j	x	0	...	0	0	0	0	0	0	1
3	9	az	t	n	f	d	x	l	e	0	...	0	0	0	0	0	0	0
4	13	az	v	n	f	d	h	d	n	0	...	0	0	0	0	0	0	0
5	18	t	b	e	c	d	g	h	s	0	...	0	0	1	0	0	0	0
6	24	al	r	e	f	d	f	h	s	0	...	0	0	0	0	0	0	0
7	25	o	l	as	f	d	f	j	a	0	...	0	0	0	0	0	0	0
8	27	w	s	as	e	d	f	i	h	0	...	1	0	0	0	0	0	0
9	30	j	b	aq	c	d	f	a	e	0	...	0	0	1	0	0	0	0

10 rows × 377 columns



```
In [189]: y.head(10)
```

```
Out[189]:
```

	y
0	130.81
1	88.53
2	76.26
3	80.62
4	78.02
5	92.93
6	128.76
7	91.91
8	108.67
9	126.99

**Written Custom class to remove 'ID' column.**

**Because , to add this object to Pipeline for future predicting values**

```
In [190]: class RemoveIDClass():  
  
    def __init__(self):  
        print('RemoveIDClass initiated. Use fit transform to remove columns')  
  
    def fit(self, X):  
        print(' Use transform to drop columns')  
        return self  
  
    def transform(self, X):  
        self.X=X  
        self.X= self.X.drop('ID', axis=1)  
        return self.X
```

**Applying RemoveIDClass to remove 'ID'**

**and apply same thing for final predicting test data i.e. new\_df here**

```
In [191]: column_selection= RemoveIDClass()
column_selection.fit(X)
X=column_selection.transform(X)
```

RemoveIDClass initiated. Use fit transform to remove columns  
Use transform to drop columns

```
In [195]: new_df= column_selection.transform(new_df)
```

```
In [196]: new_df.head()
```

Out[196]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380	X382
0	az	v	n	f	d	t	a	w	0	0	...	0	0	0	1	0	0	0
1	t	b	ai	a	d	b	g	y	0	0	...	0	0	1	0	0	0	0
2	az	v	as	f	d	a	j	j	0	0	...	0	0	0	1	0	0	0
3	az	l	n	f	d	z	l	n	0	0	...	0	0	0	1	0	0	0
4	w	s	as	c	d	y	i	m	0	0	...	1	0	0	0	0	0	0

5 rows × 376 columns



## Splitting X, y into train and test with 70, 30 %

```
In [169]: from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test=train_test_split(X,y, test_size=0.3, random_state=41)
```

## Task 1: Apply label encoder.

```
In [80]: from feature_engine import categorical_encoders as ce
```

```
# set up the encoder
```

```
encoder = ce.OneHotCategoricalEncoder(drop_last=False)
```

```
# fit the encoder
```

```
encoder.fit(X_train)
```

```
# transform the data
```

```
encoded_X_train= encoder.transform(X_train)
```

```
encoded_X_test= encoder.transform(X_test)
```

```
encoder.encoder_dict_
```

```
Out[80]: {'X0': array(['y', 'n', 'ay', 'j', 'al', 'az', 'z', 'f', 's', 'aj', 'e', 'a',
                    'x',
                    'ak', 'x', 'w', 'ap', 'o', 't', 'r', 'a', 'af', 'u', 'd', 'ai',
                    'i', 'v', 'ba', 'am', 'at', 'h', 'aq', 'l', 'as', 'm', 'c', 'b',
                    'bc', 'k', 'aw', 'q', 'au', 'ao', 'ad', 'g', 'ac', 'ab', 'aa'],
                    dtype=object),
          'X1': array(['aa', 'e', 'b', 'r', 's', 'c', 'o', 'l', 'v', 'y', 'z', 'u',
                    'i',
                    'a', 'm', 'd', 'f', 'h', 'w', 't', 'n', 'p', 'j', 'k', 'g', 'ab',
                    'q'], dtype=object),
          'X2': array(['q', 'as', 'f', 'ay', 'e', 'm', 'ai', 'aq', 'ae', 'r', 'a', 'a',
                    'h',
                    'i', 'al', 'k', 'ak', 'n', 's', 'aw', 't', 'b', 'd', 'y', 'ap',
                    'x', 'g', 'h', 'ao', 'au', 'z', 'ac', 'at', 'ag', 'an', 'av', 'am',
                    'af', 'aa', 'o', 'c', 'j', 'p', 'ar', 'l'], dtype=object),
          'X3': array(['c', 'f', 'd', 'a', 'b', 'g', 'e'], dtype=object),
          'X4': array(['d'], dtype=object),
          'X5': array(['n', 'q', 'w', 's', 'j', 'af', 'p', 'l', 'k', 'c', 'd', 'i',
                    'r',
                    'ab', 'ad', 'ag', 'm', 'ae', 'ac', 'aa', 'ah', 'v', 'o', 'f', 'g',
                    'x', 'u'], dtype=object),
          'X6': array(['j', 'a', 'h', 'd', 'g', 'i', 'l', 'k', 'c', 'f', 'e', 'b'],
                    dtype=object),
          'X8': array(['h', 'm', 'v', 'n', 'g', 'o', 'p', 'e', 'r', 'u', 'q', 'j',
                    'i',
                    'w', 'c', 'a', 'f', 't', 'y', 'k', 'b', 'x', 's', 'l', 'd'],
                    dtype=object)}
```

## Applying same label encoder to new\_Df for final predicting values

## and visualizing data using head()

```
In [170]: encoded_new_df= encoder.transform(new_df)
```

In [171]: `encoded_new_df.head(10)`

Out[171]:

	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	...	X8_a	X8_f	X8_t	X8_y	X8_k	X8_
0	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	
1	0	0	0	0	0	0	0	0	0	1	...	0	0	0	1	0	
2	0	0	0	0	1	0	0	0	0	0	...	0	0	0	0	0	
3	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	
4	0	0	0	0	1	0	0	0	0	0	...	0	0	0	0	0	
5	0	0	0	0	0	0	0	0	0	1	...	0	0	0	0	0	
6	0	0	0	0	0	0	0	0	0	0	...	0	0	0	1	0	
7	0	0	0	1	0	0	0	0	0	0	...	1	0	0	0	0	
8	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0	
9	0	0	0	0	1	0	0	0	0	0	...	0	0	0	0	0	

10 rows × 558 columns



In [81]: `X_train.head()`

Out[81]:

	X0	X1	X2	X3	X4	X5	X6	X8	X10	X11	...	X375	X376	X377	X378	X379	X380	X
2468	y	aa	q	c	d	n	j	h	0	0	...	1	0	0	0	0	0	
3023	n	e	as	c	d	q	j	m	0	0	...	0	0	1	0	0	0	
3925	y	b	f	c	d	w	a	v	0	0	...	0	0	1	0	0	0	
3999	ay	aa	as	c	d	w	h	n	0	0	...	1	0	0	0	0	0	
3196	j	aa	ay	c	d	s	j	g	1	0	...	1	0	0	0	0	0	

5 rows × 376 columns




In [82]: `encoded_X_train.head()`

Out[82]:

	X10	X11	X12	X13	X14	X15	X16	X17	X18	X19	...	X8_a	X8_f	X8_t	X8_y	X8_k
2468	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3023	0	0	0	0	1	0	0	0	0	0	...	0	0	0	0	0
3925	0	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0
3999	0	0	0	0	1	0	0	0	0	0	...	0	0	0	0	0
3196	1	0	0	0	0	0	0	0	0	0	...	0	0	0	0	0

5 rows × 558 columns



## Task 2: Missing values handling : Check for null and unique values for test and train sets

```
In [83]: from sklearn.impute import SimpleImputer
imputer= SimpleImputer()
imputer.fit(encoded_X_train)
missingvalue_encoded_X_train = imputer.transform(encoded_X_train)
missingvalue__encoded_X_test= imputer.transform(encoded_X_test)
```

## Applying same missing values to new\_Df for final predicting values

```
In [172]: missingvalue__encoded_new_df = imputer.transform(encoded_new_df)
```

```
In [84]: missingvalue_encoded_X_train.shape
```

```
Out[84]: (2946, 558)
```

```
In [85]: missingvalue__encoded_X_test.shape
```

```
Out[85]: (1263, 558)
```

```
In [86]: type(missingvalue__encoded_X_test)
```

```
Out[86]: numpy.ndarray
```

## Task 3: If for any column(s), the variance is equal to zero, then you need to remove those variable(s)

### Applying Varaince on data sets and visualizing using head()

```
In [87]: from sklearn.feature_selection import VarianceThreshold  
variance = VarianceThreshold()  
  
variance.fit(missingvalue_encoded_X_train)  
variance.transform(missingvalue_encoded_X_train)  
variance.transform(missingvalue__encoded_X_test)  
variance.variances_
```



```
Out[87]: array([0.01173937, 0.          , 0.06881193, 0.05855662, 0.24506286,
0.00067843, 0.00169434, 0.00741199, 0.00774624, 0.08874841,
0.12657157, 0.00270817, 0.07709534, 0.0199518 , 0.00169434,
0.0054016 , 0.21702065, 0.03374027, 0.04372493, 0.00405673,
0.17717697, 0.01140788, 0.00033933, 0.00640783, 0.17717697,
0.00304566, 0.17717697, 0.03215885, 0.00033933, 0.00067843,
0.01140788, 0.00033933, 0.06736417, 0.0100796 , 0.18724516,
0.24090197, 0.01173937, 0.02125242, 0.10905659, 0.16966451,
0.19687804, 0.04125089, 0.00741199, 0.04403315, 0.00640783,
0.02351964, 0.01306304, 0.24297044, 0.00101729, 0.00135593,
0.04526372, 0.0054016 , 0.01140788, 0.23506726, 0.00237046,
0.02545399, 0.00169434, 0.07111637, 0.02738003, 0.07254915,
0.09442251, 0.021577 , 0.00101729, 0.03405587, 0.04403315,
0.01173937, 0.00640783, 0.02641805, 0.0492376 , 0.17953913,
0.01635607, 0.00135593, 0.09442251, 0.24427014, 0.00135593,
0.          , 0.00640783, 0.          , 0.00674278, 0.00101729,
0.00101729, 0.          , 0.00674278, 0.00033933, 0.18397228,
0.00338291, 0.05196594, 0.00774624, 0.21321833, 0.05855662,
0.00741199, 0.16596692, 0.00169434, 0.00203251, 0.01438302,
0.          , 0.01537058, 0.03719912, 0.00067843, 0.02384261,
0.00270817, 0.02125242, 0.13012328, 0.20397766, 0.15749813,
0.04341648, 0.23481734, 0.23481734, 0.04125089, 0.00640783,
0.00270817, 0.          , 0.00304566, 0.03719912, 0.24999067,
0.039074 , 0.10931135, 0.039074 , 0.02738003, 0.21386646,
0.10777936, 0.02125242, 0.02481013, 0.04403315, 0.24166244,
0.04125089, 0.08018663, 0.04125089, 0.01569931, 0.18043728,
0.03374027, 0.1524763 , 0.00101729, 0.04125089, 0.02125242,
0.04495642, 0.16438516, 0.08130382, 0.03215885, 0.00067843,
0.16557287, 0.07426089, 0.20397766, 0.20397766, 0.18043728,
0.01372349, 0.00135593, 0.15729153, 0.0409406 , 0.21243298,
0.05825946, 0.00439329, 0.03374027, 0.00067843, 0.19640759,
0.0070775 , 0.02384261, 0.22530003, 0.0054016 , 0.01041201,
0.01733949, 0.0202773 , 0.01733949, 0.04495642, 0.24674798,
0.04495642, 0.131763 , 0.08573571, 0.09602501, 0.00405673,
0.00135593, 0.02060257, 0.24971195, 0.24601379, 0.07906575,
0.00033933, 0.24929899, 0.00270817, 0.24971195, 0.01306304,
0.01140788, 0.02897873, 0.02222547, 0.00270817, 0.00741199,
0.14024251, 0.18132967, 0.01635607, 0.00033933, 0.00033933,
0.01897391, 0.00033933, 0.06033472, 0.08983703, 0.00033933,
0.01537058, 0.00405673, 0.00169434, 0.00741199, 0.08847568,
0.0054016 , 0.00674278, 0.21577626, 0.06649275, 0.24678658,
0.00808027, 0.02125242, 0.24686309, 0.21640133, 0.08352714,
0.03215885, 0.00304566, 0.04063007, 0.0409406 , 0.00607265,
0.01668411, 0.04372493, 0.          , 0.15955138, 0.          ,
0.00033933, 0.00607265, 0.07878495, 0.00741199, 0.00338291,
0.08929318, 0.00674278, 0.00640783, 0.09442251, 0.          ,
0.2447725 , 0.18132967, 0.00135593, 0.0070775 , 0.24719596,
0.23899851, 0.00067843, 0.00135593, 0.00607265, 0.01864749,
0.06938943, 0.00033933, 0.00169434, 0.00033933, 0.00033933,
0.24274091, 0.00135593, 0.04372493, 0.037825 , 0.08683446,
0.00135593, 0.00741199, 0.          , 0.00033933, 0.00033933,
0.00237046, 0.037825 , 0.20236168, 0.01041201, 0.19858528,
0.03969712, 0.00135593, 0.00033933, 0.04372493, 0.          ,
0.00304566, 0.00439329, 0.12345148, 0.037825 , 0.16908621,
0.05557468, 0.01504163, 0.00033933, 0.          , 0.          ,
0.01140788, 0.00674278, 0.          , 0.10726685, 0.00033933,
0.00033933, 0.          , 0.00472962, 0.00472962, 0.16298902,
```

0.04618422, 0.0100796 , 0.06967783, 0.01306304, 0.04156095,  
0.00270817, 0.00941408, 0.00640783, 0.00338291, 0.23997283,  
0.00371993, 0.21150625, 0.24467387, 0.0280202 , 0.15667037,  
0.00674278, 0.00033933, 0.00067843, 0.00640783, 0.18150746,  
0.02448785, 0.0100796 , 0.24297044, 0.0054016 , 0.03215885,  
0.10956587, 0.04063007, 0.24601379, 0. , 0.05075565,  
0.00067843, 0.02287302, 0.24882462, 0.00472962, 0.11435806,  
0.24985885, 0.00674278, 0. , 0.02092761, 0.00908097,  
0.02125242, 0.07283502, 0.00908097, 0.02190135, 0.04279889,  
0. , 0.0489333 , 0.04031932, 0.22572496, 0.20934561,  
0.0489333 , 0.00203251, 0.16869952, 0.23174841, 0.14601236,  
0.00033933, 0.24697612, 0.03152467, 0.07426089, 0.03279211,  
0.24999712, 0.18622061, 0.00338291, 0.00338291, 0.00135593,  
0.04832401, 0.06063027, 0. , 0.00741199, 0.01471244,  
0.00067843, 0.01897391, 0.17570397, 0.21577626, 0.05527522,  
0.21552461, 0.02222547, 0.00974695, 0.00774624, 0.00741199,  
0.00203251, 0.00033933, 0.00135593, 0.06736417, 0.0489333 ,  
0.06269262, 0.03844996, 0.01635607, 0.04248975, 0.07709534,  
0.05196594, 0.02448785, 0.03719912, 0.0070775 , 0.00439329,  
0.07737752, 0.06620182, 0.04341648, 0.02351964, 0.06298632,  
0.06649275, 0.00270817, 0.00506572, 0.00774624, 0.00371993,  
0.01733949, 0.00841407, 0.00405673, 0.00941408, 0.00640783,  
0.00304566, 0.00573724, 0.01537058, 0.00371993, 0.00405673,  
0.00203251, 0.00841407, 0.00101729, 0.00203251, 0.00135593,  
0.00270817, 0.00405673, 0.00033933, 0.00270817, 0.00135593,  
0.00304566, 0.00033933, 0.00033933, 0.00033933, 0.00033933,  
0.15353346, 0.00841407, 0.12029244, 0.05226793, 0.12633294,  
0.02641805, 0.01766684, 0.12345148, 0.09281171, 0.00506572,  
0.00841407, 0.00974695, 0.0431078 , 0.03279211, 0.00908097,  
0.00067843, 0.00607265, 0.00741199, 0.01339338, 0.0054016 ,  
0.00506572, 0.00237046, 0.00640783, 0.00371993, 0.00203251,  
0.00067843, 0.00033933, 0.00135593, 0.23899851, 0.0199518 ,  
0.01140788, 0.01864749, 0.07709534, 0.08847568, 0.01537058,  
0.10777936, 0.03279211, 0.01140788, 0.00101729, 0.00674278,  
0.00135593, 0.0054016 , 0.06033472, 0.03374027, 0.02125242,  
0.00169434, 0.00741199, 0.0054016 , 0.00338291, 0.00203251,  
0.00169434, 0.00203251, 0.00270817, 0.00135593, 0.00405673,  
0.00101729, 0.00304566, 0.00304566, 0.00135593, 0.00405673,  
0.00135593, 0.00101729, 0.00033933, 0.00033933, 0.00033933,  
0.00033933, 0.00033933, 0.00033933, 0.00101729, 0.00033933,  
0.00033933, 0.24847619, 0.19126041, 0.06678346, 0.09469017,  
0.01372349, 0.04984551, 0.03563037, 0. , 0.0489333 ,  
0.0504525 , 0.05105857, 0.04587762, 0.02993519, 0.04495642,  
0.04771379, 0.04434114, 0.04218038, 0.0280202 , 0.04710266,  
0.04771379, 0.04771379, 0.04187078, 0.04279889, 0.04618422,  
0.04954167, 0.04832401, 0.04125089, 0.02673894, 0.02319645,  
0.0504525 , 0.00439329, 0.00135593, 0.00033933, 0.00067843,  
0.00033933, 0.18501481, 0.04862877, 0.04464889, 0.12537614,  
0.18876643, 0.10235207, 0.09735409, 0.00974695, 0.00741199,  
0.00472962, 0.00338291, 0.00674278, 0.02770023, 0.03625856,  
0.04801902, 0.05287123, 0.03025355, 0.03719912, 0.02287302,  
0.04862877, 0.05166371, 0.02673894, 0.02897873, 0.06181014,  
0.0525697 , 0.04434114, 0.02287302, 0.04832401, 0.0525697 ,  
0.02833994, 0.02641805, 0.037825 , 0.04403315, 0.02416535,  
0.05587391, 0.02092761, 0.02416535])

## Apply same variance object on new\_df

In [174]: `variance.transform(misingvalue__encoded_new_df)`

Out[174]: `array([[0., 0., 0., ..., 0., 0., 0.],  
[0., 0., 0., ..., 0., 0., 0.],  
[0., 0., 0., ..., 0., 0., 0.],  
...,  
[0., 0., 0., ..., 0., 0., 0.],  
[0., 0., 1., ..., 0., 0., 0.],  
[0., 0., 0., ..., 0., 0., 0.]])`

In [88]: `pd.DataFrame(data = misingvalue_encoded_X_train).head()`

Out[88]:

	0	1	2	3	4	5	6	7	8	9	...	548	549	550	551	552	553	554	555	!
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
1	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
2	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
3	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
4	1.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!

5 rows × 558 columns




In [89]: `pd.DataFrame(data = misingvalue__encoded_X_test).head()`

Out[89]:

	0	1	2	3	4	5	6	7	8	9	...	548	549	550	551	552	553	554	555	!
0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	!
1	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
2	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!
3	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	1.0	...	0.0	1.0	0.0	0.0	0.0	0.0	0.0	0.0	!
4	0.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	0.0	...	0.0	0.0	0.0	0.0	0.0	0.0	0.0	0.0	!

5 rows × 558 columns



## Task 4: Perform dimensionality reduction.

### Applying PCA to dimensionality reduction and print values

```
In [201]: from sklearn.decomposition import PCA
pca= PCA(n_components=0.95)

final_X_train = pca.fit_transform(missingvalue_encoded_X_train)

final_X_test = pca.transform(missingvalue__encoded_X_test)

variance_factor = pca.explained_variance_ratio_
```

## Applying PCA on new\_df

```
In [202]: final_new_df= pca.transform(missingvalue__encoded_new_df)
```

```
In [203]: varaince=np.cumsum(np.round(variance_factor, decimals=4)*100)
varaince
```


```
Out[203]: array([11.5 , 19.26, 26.52, 32.33, 37.22, 41.52, 44.88, 47.75, 50.23,
 52.34, 54.4 , 56.11, 57.62, 59.03, 60.42, 61.72, 62.92, 64.02,
 65.01, 65.94, 66.83, 67.68, 68.49, 69.27, 70. , 70.71, 71.39,
 72.04, 72.67, 73.29, 73.88, 74.44, 74.96, 75.46, 75.94, 76.39,
 76.83, 77.24, 77.65, 78.06, 78.44, 78.82, 79.19, 79.55, 79.9 ,
 80.23, 80.56, 80.87, 81.17, 81.46, 81.74, 82.02, 82.29, 82.56,
 82.82, 83.07, 83.31, 83.55, 83.79, 84.02, 84.24, 84.46, 84.68,
 84.89, 85.1 , 85.31, 85.51, 85.71, 85.91, 86.11, 86.3 , 86.49,
 86.68, 86.87, 87.06, 87.24, 87.42, 87.6 , 87.78, 87.96, 88.14,
 88.31, 88.48, 88.65, 88.82, 88.99, 89.15, 89.31, 89.47, 89.63,
 89.79, 89.95, 90.1 , 90.25, 90.4 , 90.55, 90.7 , 90.85, 90.99,
 91.13, 91.27, 91.41, 91.55, 91.68, 91.81, 91.94, 92.07, 92.2 ,
 92.33, 92.45, 92.57, 92.69, 92.81, 92.93, 93.05, 93.16, 93.27,
 93.38, 93.49, 93.6 , 93.71, 93.82, 93.92, 94.02, 94.12, 94.22,
 94.32, 94.42, 94.52, 94.61, 94.7 , 94.79, 94.88, 94.97, 95.06])
```

```
In [204]: final_X_train = pd.DataFrame(data = final_X_train)
final_X_train.head()
```

```
Out[204]:
```

	0	1	2	3	4	5	6	7	
0	0.960959	-0.408379	-1.877273	-0.254270	0.513301	0.512033	1.046224	0.654329	0.91521
1	-0.752272	1.052960	0.836741	0.077246	-1.047652	-0.059477	-1.834881	-0.033379	0.72171
2	2.968586	0.597727	-0.194642	0.238958	-0.793592	0.188629	-1.355411	0.286202	0.85684
3	-0.796384	1.620700	-1.563388	-0.505662	1.292720	1.189746	-0.407920	1.497050	1.02061
4	0.486709	0.137798	-1.145470	0.187931	0.645374	0.800998	2.484812	0.663025	-0.02371

5 rows × 135 columns




```
In [205]: final_X_test = pd.DataFrame(data = final_X_test)
          final_X_test.head()
```

Out[205]:

	0	1	2	3	4	5	6	7	
0	3.356964	0.780834	0.508713	-0.740357	-1.494637	-0.455610	-0.431170	-1.028857	-0.77251
1	1.672618	1.268216	-0.160392	1.044819	-1.166278	-0.318978	0.829785	-0.009819	-0.13971
2	-2.359431	1.657815	1.913290	-0.249068	-0.135485	0.652335	0.134101	-0.897848	0.24081
3	1.851941	-0.210400	-0.424963	-1.833808	2.352375	0.898610	0.244001	-0.549630	0.01091
4	-3.020196	-0.073882	-0.183182	-0.139958	-0.218389	-1.509716	-0.756099	0.212126	-1.15481

5 rows × 135 columns



```
In [206]: final_X_train.shape
```

Out[206]: (2946, 135)

```
In [207]: y_train.shape
```

Out[207]: (2946, 1)

## Task 5: Predict your test\_df values using XGBoost.

i. Create object XGBRegressor

ii. fit final\_X\_train and y\_train

iii. Predict for final\_X\_test

iv. Check MSE for prediction

v. Hyperparameter tuning to find optimised model

vi. Predict for the given test.csv i.e. new\_df here on optimised model

```
In [101]: from xgboost import XGBRegressor

          xgb = XGBRegressor()

          xgb.fit(final_X_train,y_train)
          y_preds = xgb.predict(final_X_test)
```

In [102]: `y_preds`

Out[102]: `array([ 90.128044, 102.15285 , 119.29104 , ..., 101.01388 , 100.81295 ,  
92.27072 ], dtype=float32)`

In [103]: `print(xgb)`

```
XGBRegressor(base_score=0.5, booster=None, colsample_bylevel=1,
              colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
              importance_type='gain', interaction_constraints=None,
              learning_rate=0.300000012, max_delta_step=0, max_depth=6,
              min_child_weight=1, missing=nan, monotone_constraints=None,
              n_estimators=100, n_jobs=0, num_parallel_tree=1,
              objective='reg:squarederror', random_state=0, reg_alpha=0,
              reg_lambda=1, scale_pos_weight=1, subsample=1, tree_method=None,
              validate_parameters=False, verbosity=None)
```

In [104]: `y_test`

Out[104]:

	<b>y</b>
<b>2107</b>	89.91
<b>3683</b>	104.87
<b>137</b>	110.02
<b>1021</b>	108.64
<b>1299</b>	104.37
...	...
<b>2242</b>	115.18
<b>1131</b>	118.42
<b>2676</b>	104.28
<b>2206</b>	88.25
<b>2481</b>	128.87

1263 rows × 1 columns

## iv. Calculating mean\_squared\_error

In [105]: `from sklearn.metrics import mean_squared_error`

`rms = np.sqrt(mean_squared_error(y_test, y_preds))`

In [106]: `rms`

Out[106]: 9.265854990325133

## v. Using GridSearchCV to hyper parameter tuning with different hyper parameters

```
In [110]: from sklearn.model_selection import GridSearchCV, RandomizedSearchCV

# Define your param grid
parameters = {'max_depth': (5,7,10), 'booster' :('gbtree','gblinear'),'learning_rate':(0.03,0.05,0.07), 'n_estimators':(100,200,250)}

#Define your Search object here:
Search_object = GridSearchCV(xgb, parameters)

#Fit your search object with your training data
Search_object.fit(final_X_train, y_train)

Search_object.best_params_
```

```
Out[110]: {'booster': 'gbtree',
           'learning_rate': 0.03,
           'max_depth': 5,
           'n_estimators': 250}
```

**Found that {'booster': 'gbtree', 'learning\_rate': 0.03, 'max\_depth': 5, 'n\_estimators': 250} are best params**

**So, Find the score**

```
In [111]: print("best XGboost regression from grid search: %f" % Search_object.best_estimator_.score(final_X_test, y_test))

best XGboost regression from grid search: 0.503886
```

**Predicting final X test values using Optimized XGBRegressor model**

```
In [112]: xgb_optimized = XGBRegressor(booster='gbtree', learning_rate= 0.03, max_depth=
5, n_estimators= 250)
print(xgb_optimized)
xgb_optimized.fit(final_X_train,y_train)

y_preds_optimized = xgb_optimized.predict(final_X_test)
```

```
XGBRegressor(base_score=None, booster='gbtree', colsample_bylevel=None,
             colsample_bynode=None, colsample_bytree=None, gamma=None,
             gpu_id=None, importance_type='gain', interaction_constraints=None,
             learning_rate=0.03, max_delta_step=None, max_depth=5,
             min_child_weight=None, missing=nan, monotone_constraints=None,
             n_estimators=250, n_jobs=None, num_parallel_tree=None,
             objective='reg:squarederror', random_state=None, reg_alpha=None,
             reg_lambda=None, scale_pos_weight=None, subsample=None,
             tree_method=None, validate_parameters=False, verbosity=None)
```

## Printing predicitons

```
In [208]: y_preds_optimized = pd.DataFrame(data = y_preds_optimized)
print(y_preds_optimized)
```

```
0
0    92.019440
1   110.998375
2   107.257530
3   110.801277
4   110.453041
...
1258 106.840363
1259 110.301445
1260 101.018021
1261  93.686005
1262  93.001305
```

```
[1263 rows x 1 columns]
```

## vi. Predicting values for final\_new\_df using oprimized trained XGBRegressor.




```
In [178]: final_new_df = pd.DataFrame(data = final_new_df)
          final_new_df.head()
```

Out[178]:

	0	1	2	3	4	5	6	7	
0	-0.410161	-3.089787	0.611238	3.306720	-0.608157	3.628391	-0.555998	0.185647	-1.50310
1	3.918156	0.636312	1.344518	-0.487245	0.875199	-0.306045	-2.128026	-1.104256	-0.99765
2	-1.293756	-0.794236	0.545889	1.919805	0.248940	2.964675	-0.889490	0.842970	-1.26485
3	-0.449481	-3.074606	0.463103	3.315172	-0.715547	3.547886	-0.692836	0.137772	-1.67249
4	-2.860020	1.425717	0.808587	-1.509947	0.334650	0.356097	-0.198072	0.152543	-0.00536

5 rows × 135 columns



```
In [179]: y_preds_new_df = xgb_optimized.predict(final_new_df)
```

## printing predictions for final new\_df

```
In [209]: y_preds_new_df = pd.DataFrame(data=y_preds_new_df)
          print(y_preds_new_df)
```

	0
0	78.897202
1	95.644554
2	79.189850
3	79.186584
4	111.400620
...	...
4204	105.252731
4205	94.326477
4206	96.404037
4207	110.042572
4208	93.147545

[4209 rows x 1 columns]

## Using pipeline for future predictions datasets and adding all objects to pipeline

```
In [210]: from sklearn.pipeline import Pipeline, make_pipeline

mercedesBenz_XGboostRegressor_pipeline = make_pipeline(column_selection,
                                                         encoder,
                                                         imputer,
                                                         variance,
                                                         pca,
                                                         xgb_optimized)
```

## Storing Pipeline object using Pickle

```
In [116]: import pickle

# Save your object
pickle.dump(mercedesBenz_XGboostRegressor_pipeline, open("mercedesBenz_XGboostR
egressor.pkl", 'wb'))
```

## this is how, we load pipeline for future predictions

```
In [118]: # Load you object
trained_xgboostRegressor = pickle.load(open("mercedesBenz_XGboostRegressor.pk
l", 'rb'))
```

```
In [119]: new_df = pd.read_csv("E:\\Simplilean\\ML\\ProjectSubmission\\MercedesBenz\\Data
sets\\new_test.csv", sep=",")
```

```
In [211]: type(trained_xgboostRegressor)
```

```
Out[211]: sklearn.pipeline.Pipeline
```

In [135]: `trained_xgboostRegressor.named_steps`

```
Out[135]: {'removeidclass': <__main__.RemoveIDClass at 0x1cacf898>,
  'onehotcategoricalencoder': OneHotCategoricalEncoder(drop_last=False, top_categories=None,
  variables=['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6',
  'X8']),
  'simpleimputer': SimpleImputer(add_indicator=False, copy=True, fill_value=None,
  missing_values=nan, strategy='mean', verbose=0),
  'variancethreshold': VarianceThreshold(threshold=0.0),
  'pca': PCA(copy=True, iterated_power='auto', n_components=0.95, random_state=None,
  svd_solver='auto', tol=0.0, whiten=False),
  'xgbregressor': XGBRegressor(base_score=0.5, booster='gbtree', colsample_bylevel=1,
  colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
  importance_type='gain', interaction_constraints=None,
  learning_rate=0.03, max_delta_step=0, max_depth=5,
  min_child_weight=1, missing=nan, monotone_constraints=None,
  n_estimators=250, n_jobs=0, num_parallel_tree=1,
  objective='reg:squarederror', random_state=0, reg_alpha=0,
  reg_lambda=1, scale_pos_weight=1, subsample=1, tree_method=None,
  validate_parameters=False, verbosity=None)}
```

**Finally we apply loaded object to on new future predictions datasets**

In [ ]: `trained_xgboostRegressor.predict(new_df)`

**End of Project**

**Thank you.**